



Flight Booking Price Prediction



Agenda

01 Importing the Libraries

02 Loading the Data

03 Data Visualization

04 One Hot Encoding

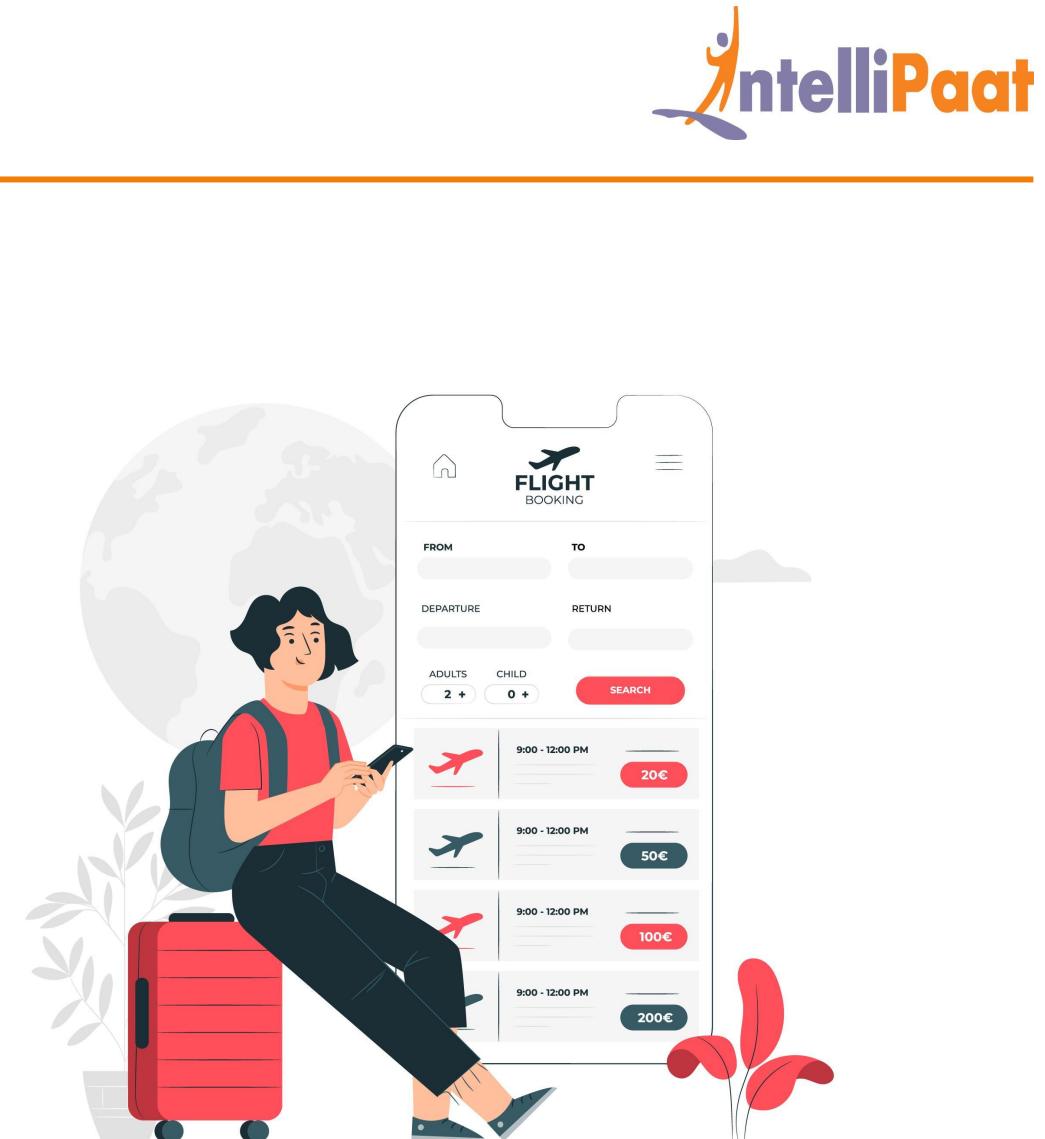
05 Feature Selection

06 Implementing ML Algorithms

Problem Statement

Problem Statement

The objective is to analyze the flight booking dataset obtained from a platform which is used to book flight tickets. A thorough study of the data will aid in the discovery of valuable insights that will be of enormous value to passengers. Apply EDA, statistical methods and Machine learning algorithms in order to get meaningful information from it.



Flight booking price prediction dataset contains around 3 lacs records with 11 attributes .



Dataset Information

| Attributes | Description |
|------------------|---|
| Airline | Name of the airline company |
| Flight | Plane's flight code |
| Source City | City from which the flight takes off |
| Departure Time | Time of Departure |
| Stops | Number of stops between the source and destination cities |
| Arrival Time | Time of Arrival |
| Destination City | City where the flight will land |
| Class | Contains information on seat class |
| Duration | Overall amount of time taken to travel between cities in hours. |
| Days left | Subtracting the trip date by the booking date. |
| Price | Ticket price |



Importing the Libraries

Importing the Libraries

We start off this project by importing all the necessary libraries that will be required for the process.

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

Loading the Data

Loading the Data

Loading the data and removing unnecessary column from the dataframe

```
import pandas as pd
df=pd.read_csv("Flight_Booking.csv")
df=df.drop(columns=["Unnamed: 0"])
df.head()
```

| | airline | flight | source_city | departure_time | stops | arrival_time | destination_city | class | duration | days_left | price |
|---|----------|---------|-------------|----------------|-------|---------------|------------------|---------|----------|-----------|-------|
| 0 | SpiceJet | SG-8709 | Delhi | Evening | zero | Night | Mumbai | Economy | 2.17 | 1 | 5953 |
| 1 | SpiceJet | SG-8157 | Delhi | Early_Morning | zero | Morning | Mumbai | Economy | 2.33 | 1 | 5953 |
| 2 | AirAsia | I5-764 | Delhi | Early_Morning | zero | Early_Morning | Mumbai | Economy | 2.17 | 1 | 5956 |
| 3 | Vistara | UK-995 | Delhi | Morning | zero | Afternoon | Mumbai | Economy | 2.25 | 1 | 5955 |
| 4 | Vistara | UK-963 | Delhi | Morning | zero | Morning | Mumbai | Economy | 2.33 | 1 | 5955 |

Loading the Data

```
df.shape  
df.info()  
df.describe()
```

(300153, 11)

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 300153 entries, 0 to 300152  
Data columns (total 11 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   airline          300153 non-null   object    
 1   flight           300153 non-null   object    
 2   source_city      300153 non-null   object    
 3   departure_time   300153 non-null   object    
 4   stops            300153 non-null   object    
 5   arrival_time     300153 non-null   object    
 6   destination_city 300153 non-null   object    
 7   class             300153 non-null   object    
 8   duration          300153 non-null   float64  
 9   days_left         300153 non-null   int64     
 10  price             300153 non-null   int64     
dtypes: float64(1), int64(2), object(8)  
memory usage: 25.2+ MB
```

Checking the shape of a dataframe and datatypes of all columns along with calculating the statistical data.

| | duration | days_left | price |
|--------------|---------------|---------------|---------------|
| count | 300153.000000 | 300153.000000 | 300153.000000 |
| mean | 12.221021 | 26.004751 | 20889.660523 |
| std | 7.191997 | 13.561004 | 22697.767366 |
| min | 0.830000 | 1.000000 | 1105.000000 |
| 25% | 6.830000 | 15.000000 | 4783.000000 |
| 50% | 11.250000 | 26.000000 | 7425.000000 |
| 75% | 16.170000 | 38.000000 | 42521.000000 |
| max | 49.830000 | 49.000000 | 123071.000000 |

Checking out the missing values in a dataframe

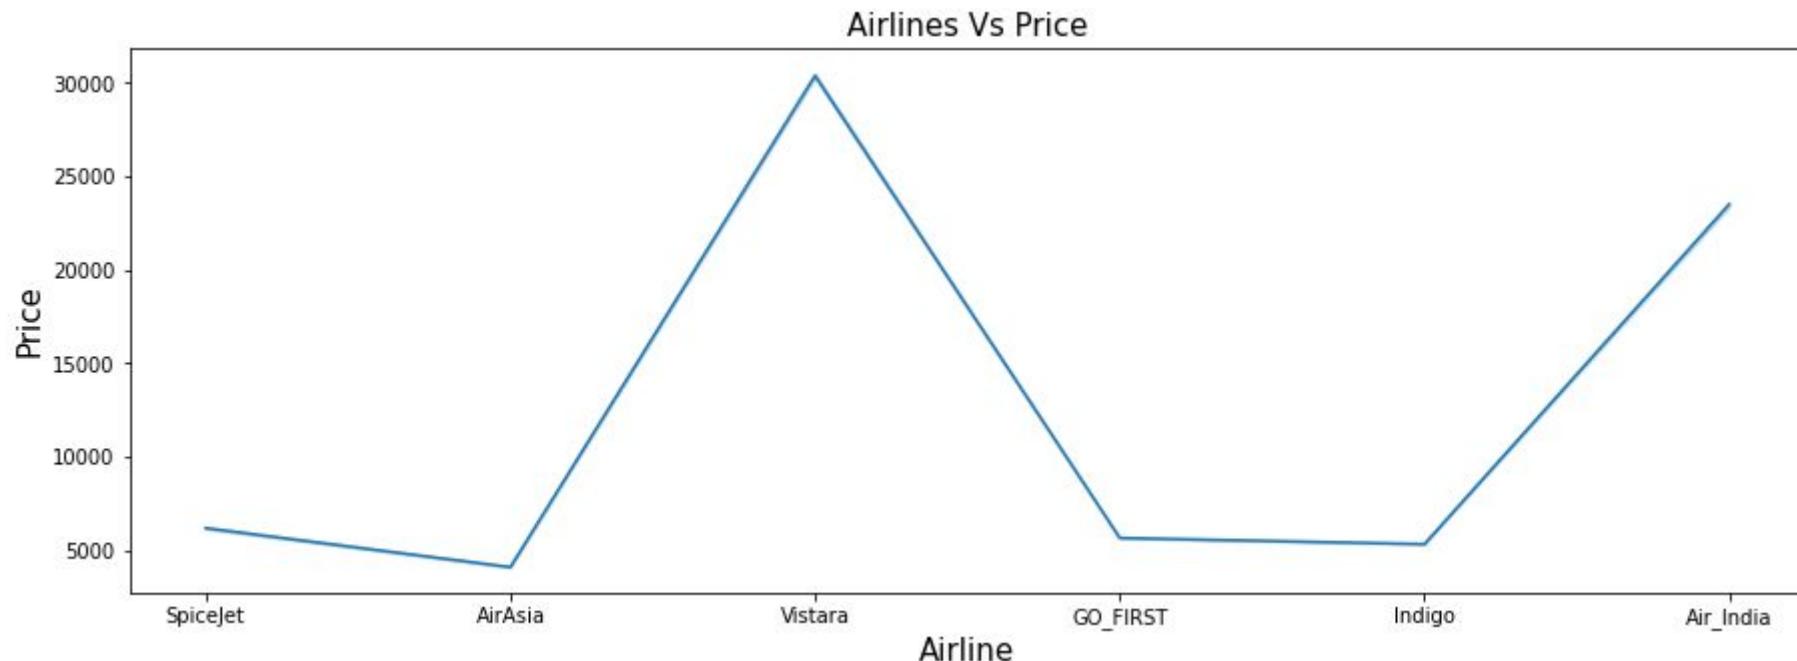
```
df.isnull().sum()
```

```
airline      0
flight       0
source_city   0
departure_time 0
stops        0
arrival_time  0
destination_city 0
class         0
duration      0
days_left     0
price         0
dtype: int64
```

Data Visualization

```
plt.figure(figsize=(15,5))
sns.lineplot(x=df['airline'],y=df['price'])
plt.title('Airlines Vs Price',fontsize=15)
plt.xlabel('Airline',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```

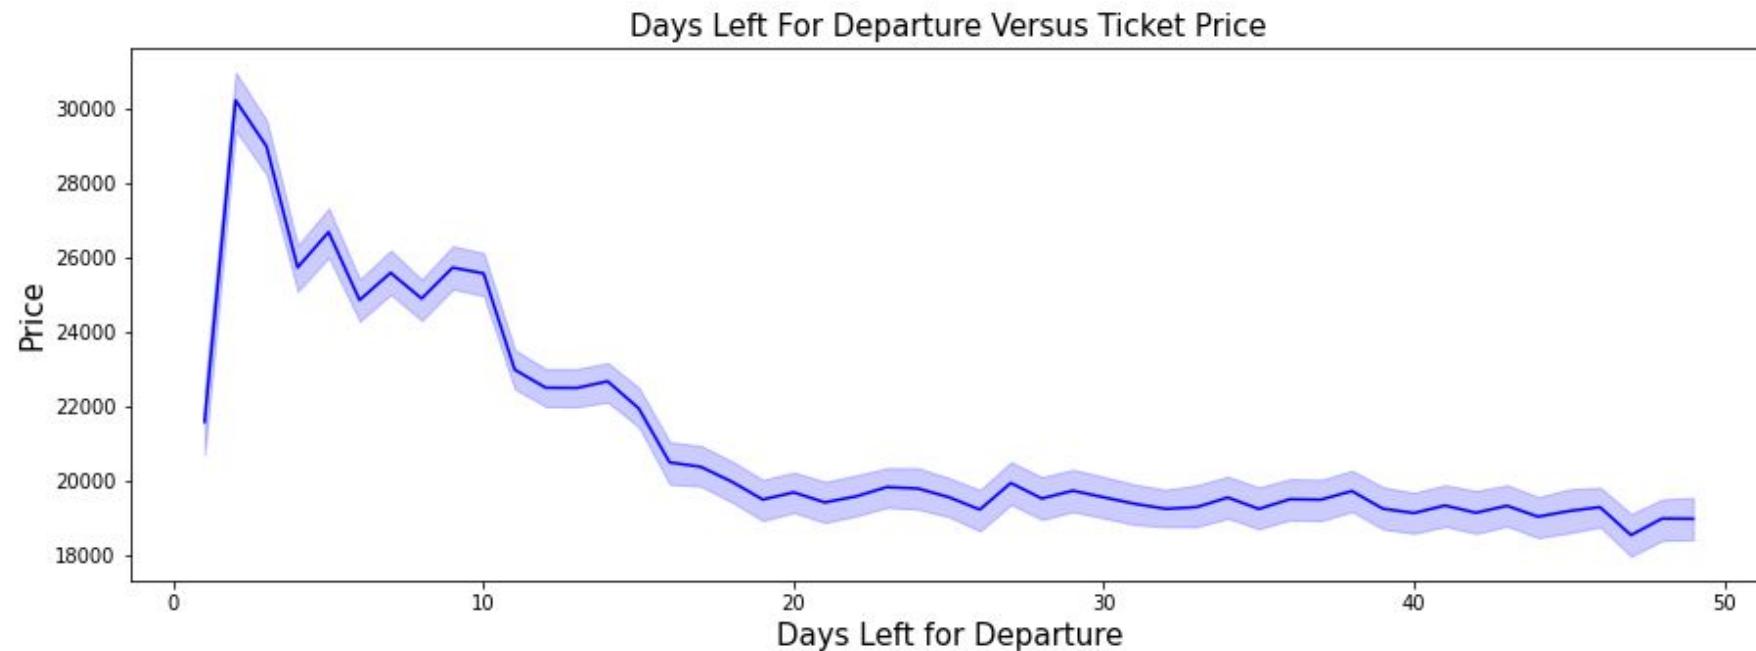
There is a variation in price with different airlines



Data Visualization

```
plt.figure(figsize=(15,5))
sns.lineplot(data=df,x='days_left',y='price',color='blue')
plt.title('Days Left For Departure Versus Ticket Price',fontsize=15)
plt.xlabel('Days Left for Departure',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```

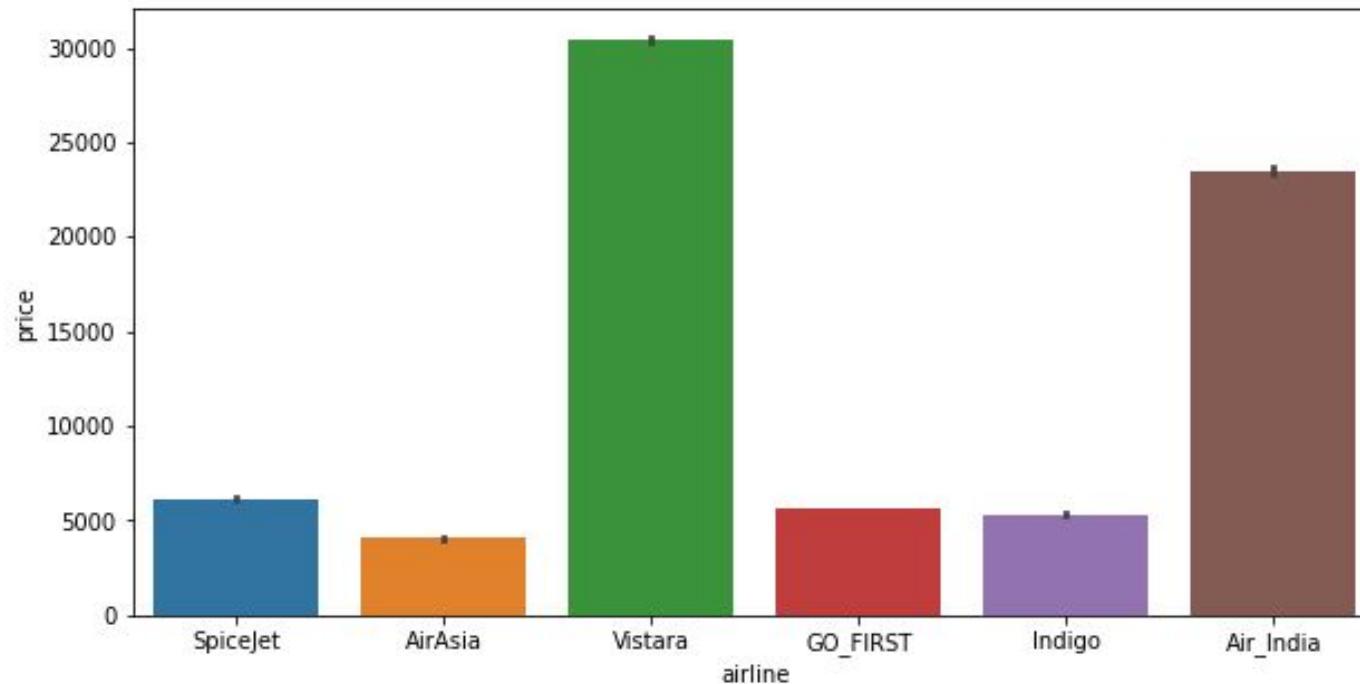
The price of the ticket increases as the days left for departure decreases



Data Visualization

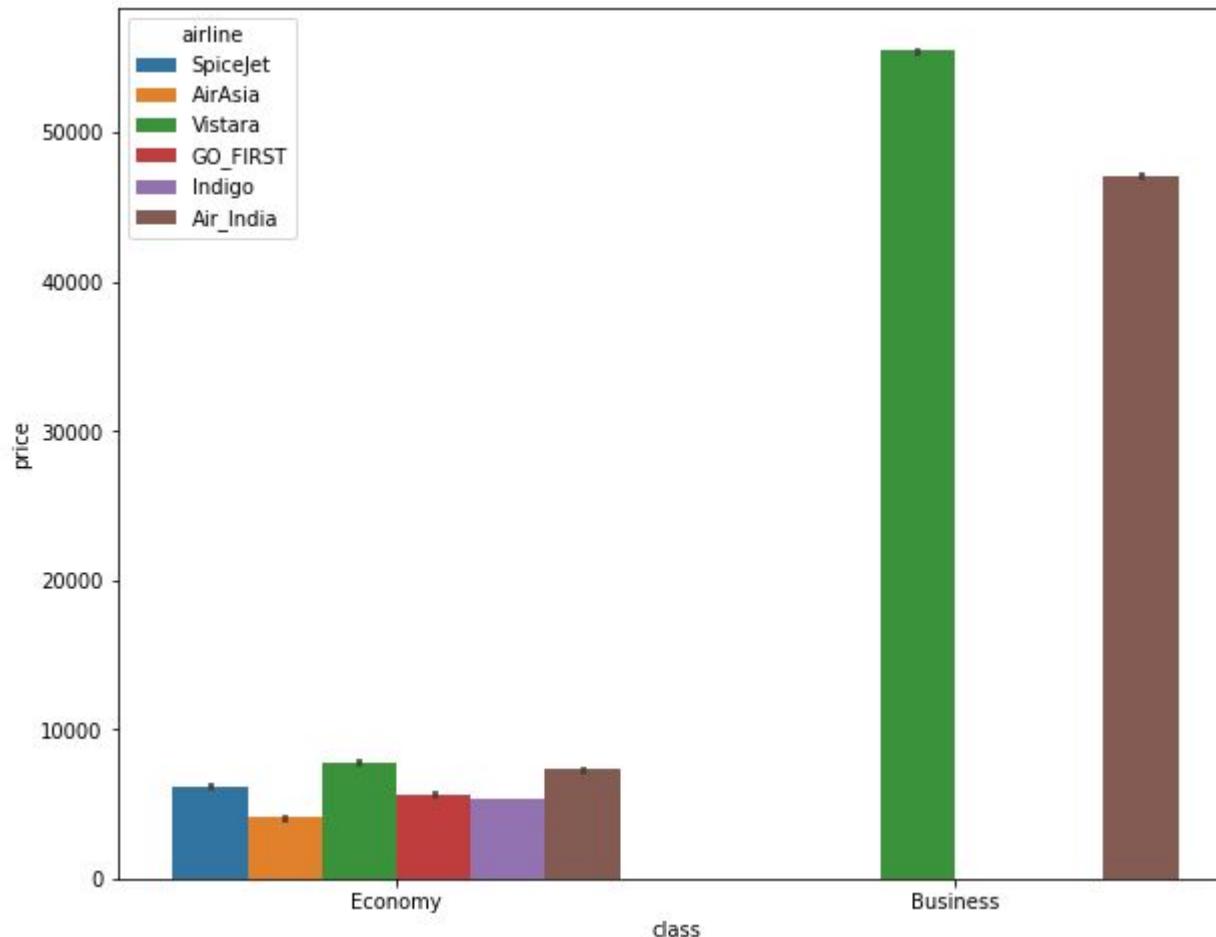
```
plt.figure(figsize=(10,5));
sns.barplot(x='airline',y='price',data=df)
```

Price range of all the flights



Data Visualization

```
plt.figure(figsize=(10,8));
sns.barplot(x='class',y='price',data=df,hue='airline')
```

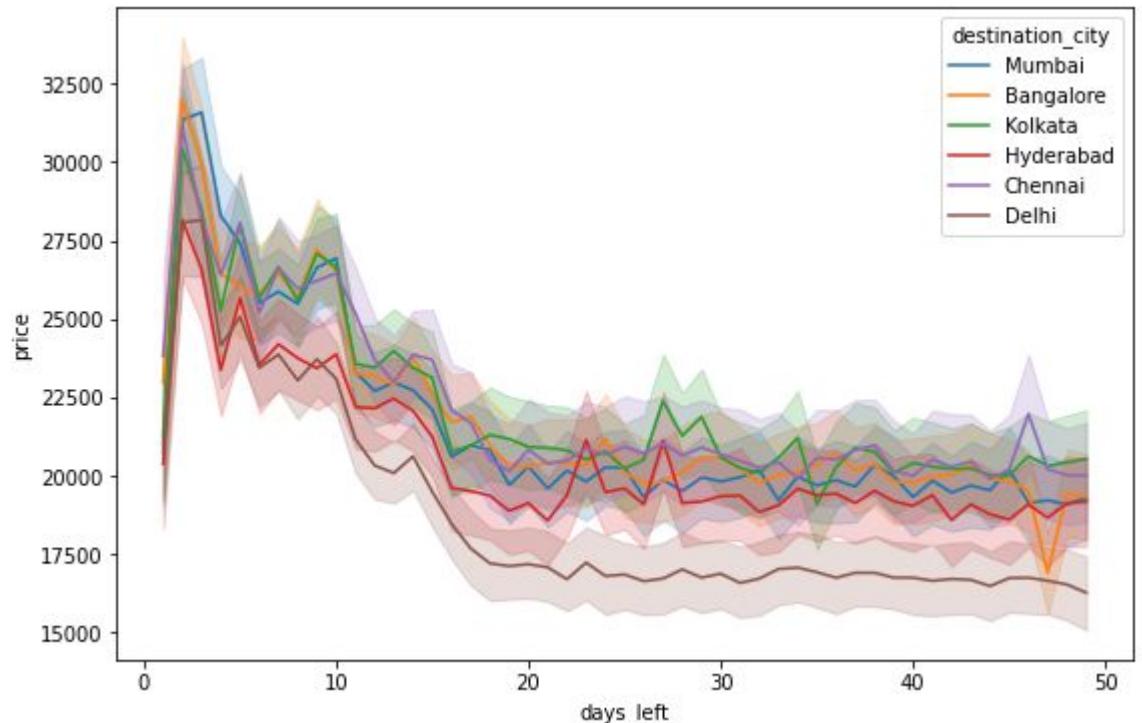
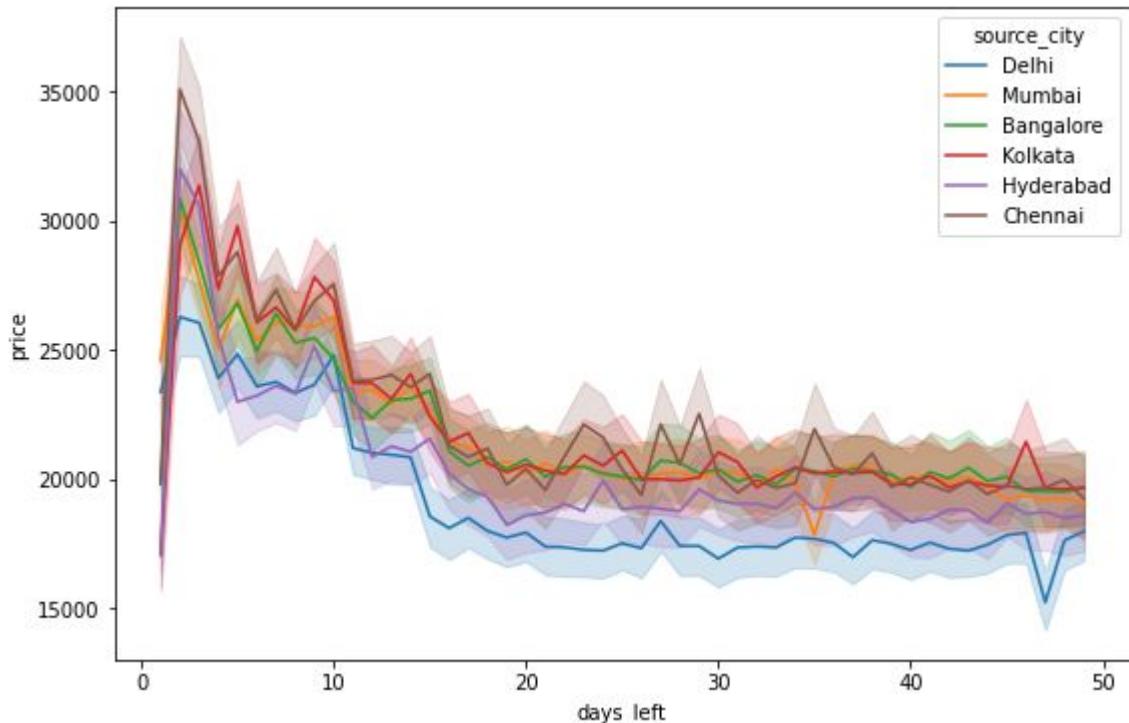


Range of price of all the flights of Economy and Business class

Data Visualization

```
fig,ax=plt.subplots(1,2,figsize=(20,6))
sns.lineplot(x='days_left',y='price',data=df,hue='source_city',ax=ax[0])
sns.lineplot(x='days_left',y='price',data=df,hue='destination_city',ax=ax[1])
plt.show()
```

Range of price of flights with source and destination city according to the days left



Data Visualization

```
plt.figure(figsize=(15,23))

plt.subplot(4, 2, 1)
sns.countplot(x=df["airline"], data=df)
plt.title("Frequency of Airline")

plt.subplot(4, 2, 2)
sns.countplot(x=df["source_city"], data=df)
plt.title("Frequency of Source City")

plt.subplot(4, 2, 3)
sns.countplot(x=df["departure_time"], data=df)
plt.title("Frequency of Departure Time")

plt.subplot(4, 2, 4)
sns.countplot(x=df["stops"], data=df)
plt.title("Frequency of Stops")

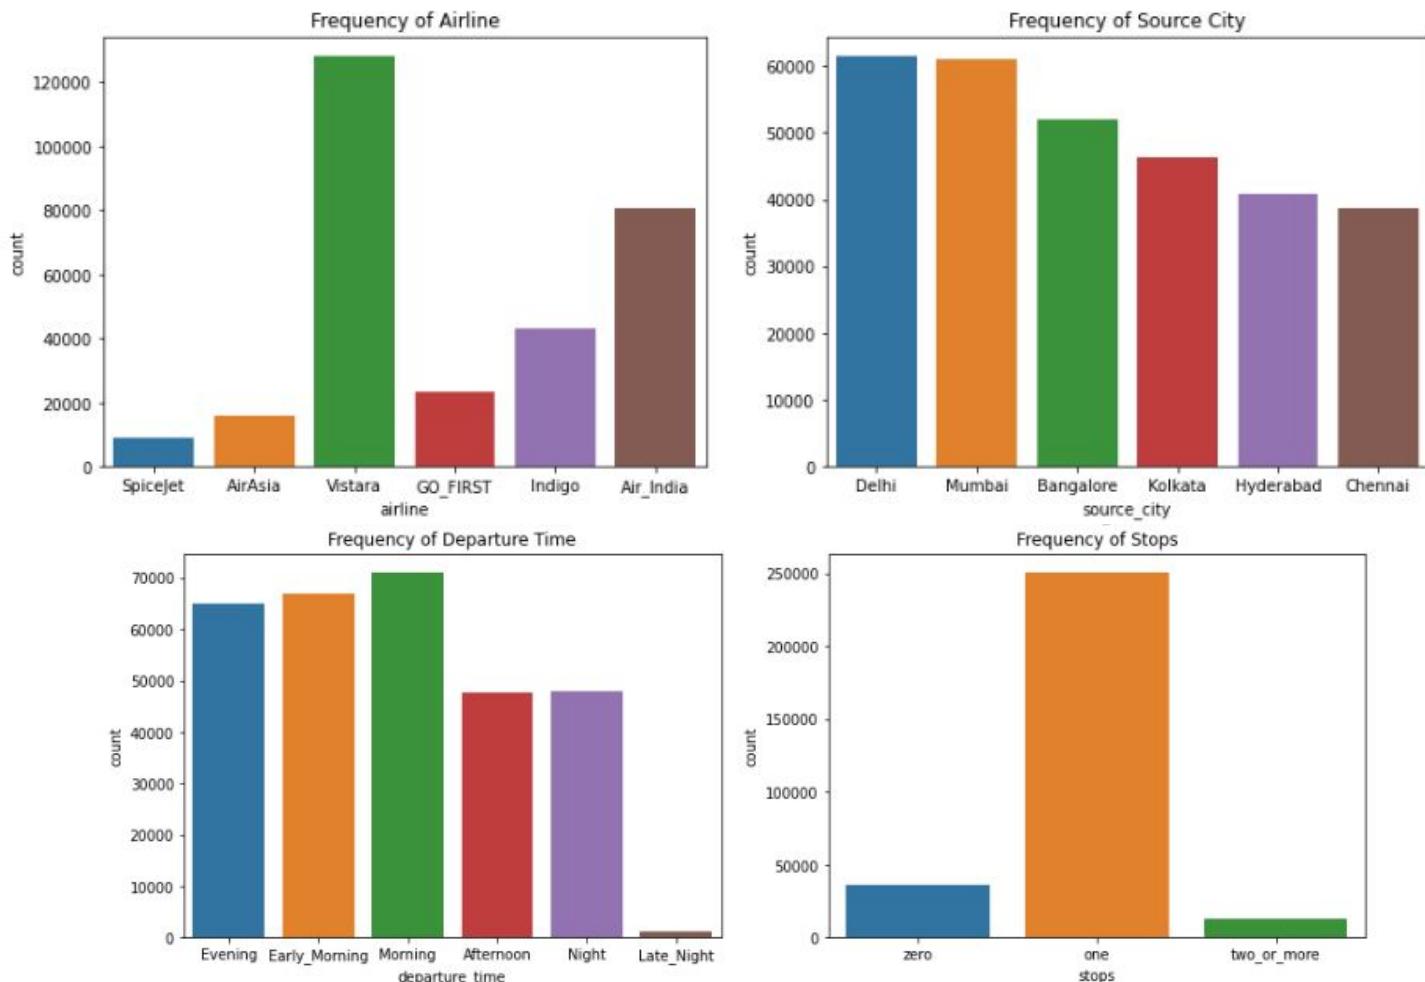
plt.subplot(4, 2, 5)
sns.countplot(x=df["arrival_time"], data=df)
plt.title("Frequency of Arrival Time")

plt.subplot(4, 2, 6)
sns.countplot(x=df["destination_city"], data=df)
plt.title("Frequency of Destination City")

plt.subplot(4, 2, 7)
sns.countplot(x=df["class"], data=df)
plt.title("Class Frequency")

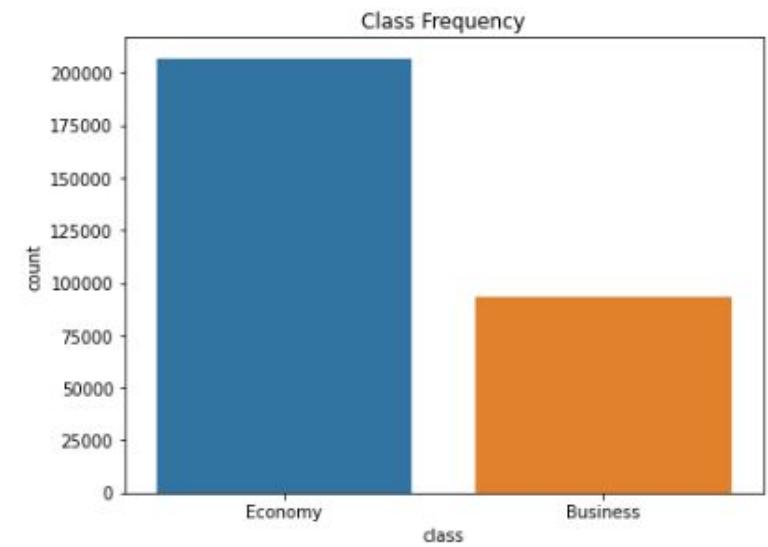
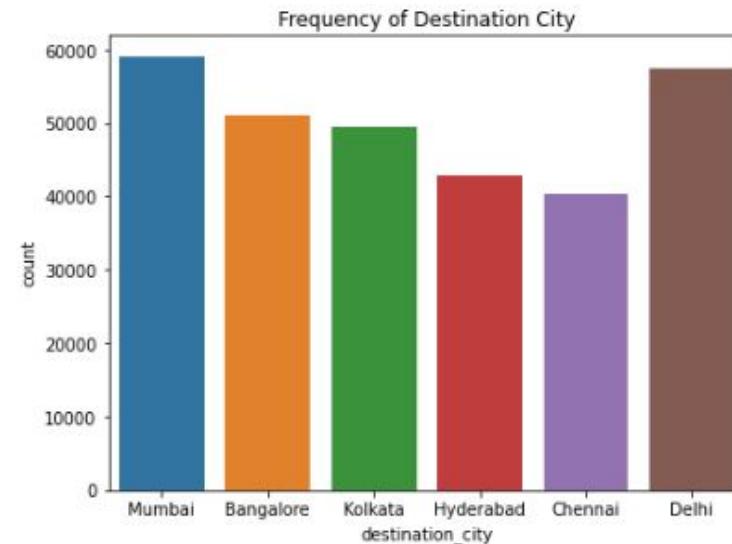
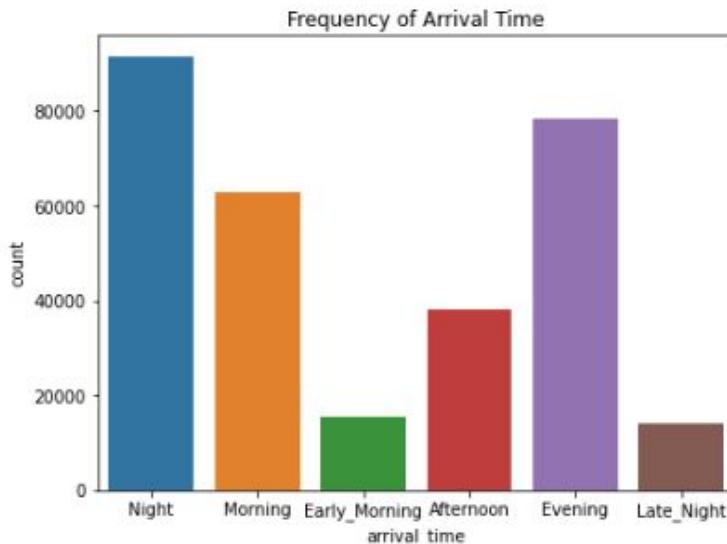
plt.show()
```

Visualization of categorical features with countplot



Data Visualization

Visualization of categorical features with countplot



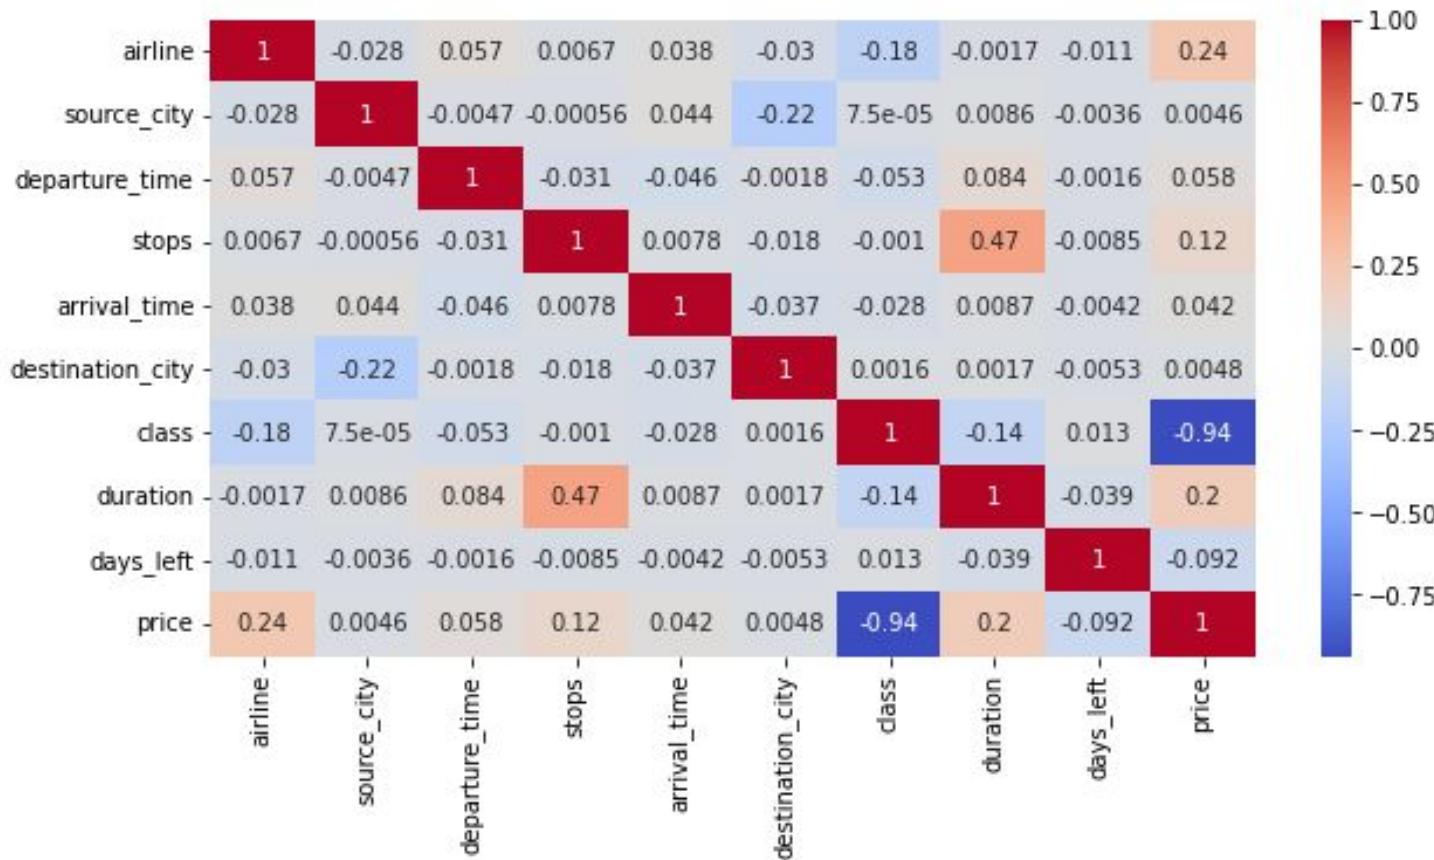
Performing One Hot Encoding for categorical features of a dataframe

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df["airline"]=le.fit_transform(df["airline"])
df["source_city"]=le.fit_transform(df["source_city"])
df["departure_time"]=le.fit_transform(df["departure_time"])
df["stops"]=le.fit_transform(df["stops"])
df["arrival_time"]=le.fit_transform(df["arrival_time"])
df["destination_city"]=le.fit_transform(df["destination_city"])
df["class"]=le.fit_transform(df["class"])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300153 entries, 0 to 300152
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   airline          300153 non-null   int64  
 1   flight           300153 non-null   object  
 2   source_city      300153 non-null   int64  
 3   departure_time   300153 non-null   int64  
 4   stops            300153 non-null   int64  
 5   arrival_time     300153 non-null   int64  
 6   destination_city 300153 non-null   int64  
 7   class             300153 non-null   int64  
 8   duration          300153 non-null   float64 
 9   days_left         300153 non-null   int64  
 10  price             300153 non-null   int64  
dtypes: float64(1), int64(9), object(1)
memory usage: 25.2+ MB
```

Feature Selection

```
plt.figure(figsize=(10,5))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
plt.show()
```



Plotting the correlation graph to see the correlation between features and dependent variable.

Feature Selection

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
col_list = []
for col in df.columns:
    if ((df[col].dtype != 'object') & (col != 'price') ):
        col_list.append(col)

X = df[col_list]
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                  for i in range(len(X.columns))]
print(vif_data)
```

| feature | VIF |
|------------------|----------|
| airline | 3.461766 |
| source_city | 2.933064 |
| departure_time | 2.746367 |
| stops | 7.464236 |
| arrival_time | 3.684695 |
| destination_city | 2.893218 |
| class | 2.917521 |
| duration | 5.037943 |
| days_left | 4.035735 |

Selecting the features using VIF. VIF should be less than 5. So drop the stops feature.

Feature Selection

```
df=df.drop(columns=["stops"])

from statsmodels.stats.outliers_influence import variance_inflation_factor
col_list = []
for col in df.columns:
    if ((df[col].dtype != 'object') & (col != 'price') ):
        col_list.append(col)

X = df[col_list]
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                  for i in range(len(X.columns))]
print(vif_data)
```

Dropping the stops column.
All features are having VIF less than 5.

| feature | VIF |
|------------------|----------|
| airline | 3.370020 |
| source_city | 2.895803 |
| departure_time | 2.746255 |
| arrival_time | 3.632792 |
| destination_city | 2.857808 |
| class | 2.776721 |
| duration | 3.429344 |
| days_left | 3.950132 |

Linear Regression

Applying standardization and implementing Linear Regression Model to predict the price of a flight.

```
X = df.drop(columns=["price"])
y = df['price']
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform(x_test)
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
y_pred=lr.predict(x_test)
difference=pd.DataFrame(np.c_[y_test,y_pred],columns=["Actual_Value","Predicted_Value"])
difference
```

| | Actual_Value | Predicted_Value |
|-------|--------------|-----------------|
| 0 | 7366.0 | 4673.755319 |
| 1 | 64831.0 | 51713.744720 |
| 2 | 6195.0 | 6610.897658 |
| 3 | 60160.0 | 55489.844234 |
| 4 | 6578.0 | 5120.342596 |
| ... | ... | ... |
| 60026 | 5026.0 | 4960.777767 |
| 60027 | 3001.0 | 4693.865426 |
| 60028 | 6734.0 | 4974.962678 |
| 60029 | 5082.0 | 2729.650066 |
| 60030 | 66465.0 | 59638.748598 |

Linear Regression

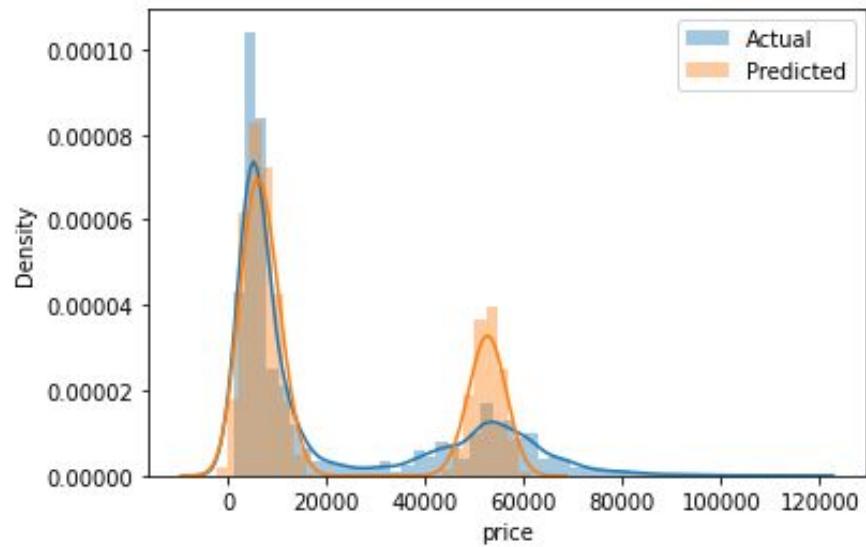
Calculating r2 score, MAE, MAPE, MSE, RMSE. Root Mean square error(RMSE) of the Linear regression model is 7259.93 and Mean absolute percentage error(MAPE) is 34 percent. Lower the RMSE and MAPE better the model.

```
from sklearn.metrics import r2_score
r2_score(y_test,y_pred)
from sklearn import metrics
mean_abs_error= metrics.mean_absolute_error(y_test,y_pred)
mean_abs_error
from sklearn.metrics import mean_absolute_percentage_error
mean_absolute_percentage_error(y_test, y_pred)
mean_sq_error=metrics.mean_squared_error(y_test,y_pred)
mean_sq_error
root_mean_sq_error = np.sqrt(metrics.mean_squared_error(y_test,y_pred))
root_mean_sq_error
```

| |
|--------------------|
| 0.897752737512321 |
| 4468.426673542113 |
| 0.3476580461068184 |
| 52706651.33334208 |
| 7259.934664536733 |

Linear Regression

```
sns.distplot(y_test,label="Actual")
sns.distplot(y_pred,label="Predicted")
plt.legend()
```



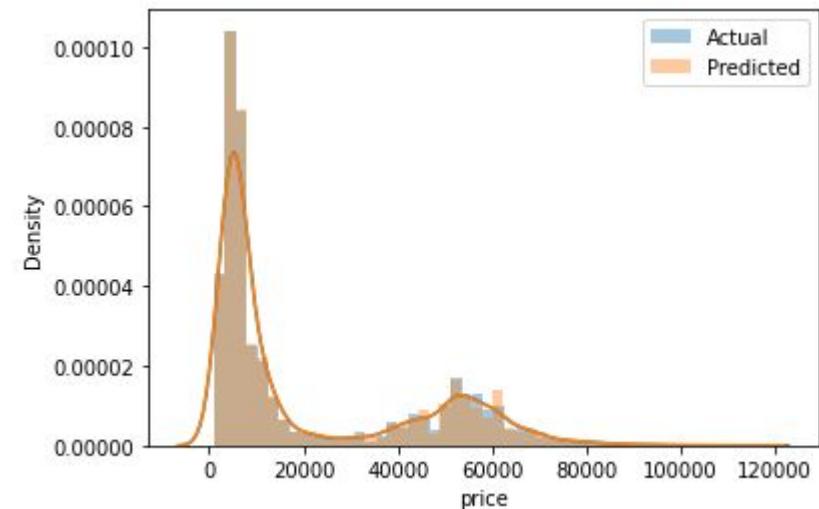
Plotting the graph of actual and predicted price of flight

Decision Tree Regressor

```
from sklearn.tree import DecisionTreeRegressor  
dt=DecisionTreeRegressor()  
dt.fit(x_train,y_train)  
y_pred=dt.predict(x_test)  
r2_score(y_test,y_pred)  
mean_abs_error= metrics.mean_absolute_error(y_test,y_pred)  
mean_abs_error  
from sklearn.metrics import mean_absolute_percentage_error  
mean_absolute_percentage_error(y_test, y_pred)  
mean_sq_error=metrics.mean_squared_error(y_test,y_pred)  
mean_sq_error  
root_mean_sq_error = np.sqrt(metrics.mean_squared_error(y_test,y_pred))  
root_mean_sq_error
```

0.9745774442285287
1219.455742310917
0.07732296917115203
13104876.849009493
3620.0658625237047

Mean absolute percentage error is 7.7 percent and RMSE is 3620 which is less than the linear regression model



Random Forest Regressor

```
from sklearn.ensemble import RandomForestRegressor
rfr=RandomForestRegressor()
rfr.fit(x_train,y_train)
y_pred=rfr.predict(x_test)
r2_score(y_test,y_pred)
mean_abs_error= metrics.mean_absolute_error(y_test,y_pred)
mean_abs_error
from sklearn.metrics import mean_absolute_percentage_error
mean_absolute_percentage_error(y_test, y_pred)
mean_sq_error=metrics.mean_squared_error(y_test,y_pred)
mean_sq_error
root_mean_sq_error = np.sqrt(metrics.mean_squared_error(y_test,y_pred))
root_mean_sq_error
```

```
sns.distplot(y_test,label="Actual")
sns.distplot(y_pred,label="Predicted")
plt.legend()
```

0.9845246238799552
1122.6731295238862
0.07319114674216119
7977282.066694117
2824.4082684155487

Mean absolute percentage error is 7.3 percent and RMSE is 2824 which is less than the linear regression and decision tree model

